

LAB 1

Implementation and Performance Analysis of Q-Learning Using an Epsilon-Greedy Policy

This lab uses Gymnasium's deterministic `CliffWalking-v1` environment instead of an Atari environment. Its discrete state/action spaces make a tabular Q-learning implementation easy to inspect.

Aim and tasks

Train a Q-learning agent with an epsilon-greedy policy using **500, 1,000, 2,000, 5,000, and 10,000 episodes**. Epsilon starts high and decays during training. The notebook records every episode's reward, plots cumulative and latest-100 average reward, evaluates each learned policy for 100 greedy episodes, compares success rates, and visualizes the environment and extracted policy.

```
In [1]: import gymnasium as gym
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd

EPISODE_CONFIGS = [500, 1_000, 2_000, 5_000, 10_000]
TEST_EPISODES = 100
MAX_STEPS = 100
SEED = 42
LEARNING_RATE = 0.8
DISCOUNT_FACTOR = 0.95
EPSILON_START = 1.0
EPSILON_MIN = 0.01
EPSILON_DECAY = 0.001
```

```
In [2]: def train_q_learning(number_of_episodes, seed=SEED):
    env = gym.make("CliffWalking-v1")
    rng = np.random.default_rng(seed)
    q_table = np.zeros((env.observation_space.n, env.action_space.n))
    rewards, epsilon = [], EPSILON_START

    for episode in range(number_of_episodes):
        state, _ = env.reset(seed=seed + episode)
        total_reward = 0.0
        for _ in range(MAX_STEPS):
            if rng.random() < epsilon:
                action = int(rng.integers(env.action_space.n))
            else:
                action = int(np.argmax(q_table[state]))
            next_state, reward, terminated, truncated, _ = env.step(action)
```

```

        target = reward if terminated else reward + DISCOUNT_FACTOR * np.max(q_
        q_table[state, action] += LEARNING_RATE * (target - q_table[state, acti
        state, total_reward = next_state, total_reward + reward
        if terminated or truncated:
            break
        rewards.append(total_reward)
        epsilon = max(EPSILON_MIN, epsilon * np.exp(-EPSILON_DECAY))
    env.close()
    return q_table, np.asarray(rewards)

def evaluate_policy(q_table, number_of_episodes=TEST_EPISODES, seed=SEED):
    env = gym.make("CliffWalking-v1")
    successes = 0
    for episode in range(number_of_episodes):
        state, _ = env.reset(seed=seed + episode)
        for _ in range(MAX_STEPS):
            state, _, terminated, truncated, _ = env.step(int(np.argmax(q_table[sta
            if terminated:
                successes += 1
            if terminated or truncated:
                break
    env.close()
    return 100 * successes / number_of_episodes

```

```

In [3]: results, summary = {}, []
        for episodes in EPISODE_CONFIGS:
            q_table, rewards = train_q_learning(episodes)
            results[episodes] = {"q_table": q_table, "rewards": rewards}
            summary.append({
                "Training episodes": episodes,
                "Final cumulative reward": rewards.sum(),
                "Average reward (last 100)": rewards[-100:].mean(),
                "Test success rate (%)": evaluate_policy(q_table),
            })
        summary_df = pd.DataFrame(summary)
        summary_df

```

```

Out[3]:

```

	Training episodes	Final cumulative reward	Average reward (last 100)	Test success rate (%)
0	500	-373334.0	-432.41	100.0
1	1000	-495121.0	-170.93	100.0
2	2000	-594297.0	-61.61	100.0
3	5000	-673884.0	-17.32	100.0
4	10000	-753359.0	-17.37	100.0

```

In [4]: fig, axes = plt.subplots(1, 2, figsize=(15, 5))
        for episodes, result in results.items():
            rewards = result["rewards"]
            axes[0].plot(np.cumsum(rewards), label=f"{episodes:,} episodes")



```

```

    axes[1].plot(pd.Series(rewards).rolling(100, min_periods=1).mean(), label=f"{ep
axes[0].set(title="Cumulative reward", xlabel="Training episode", ylabel="Cumulative
axes[1].set(title="Average reward over latest 100 episodes", xlabel="Training episo
for ax in axes:
    ax.grid(alpha=0.3)
    ax.legend()
plt.tight_layout()
plt.show()

plt.figure(figsize=(8, 4))
plt.bar(summary_df["Training episodes"].astype(str), summary_df["Test success rate
plt.title("Policy success rate after training")
plt.xlabel("Training episodes")
plt.ylabel("Success rate (%)")
plt.ylim(0, 105)
plt.grid(axis="y", alpha=0.3)
plt.show()

```




Environment and extracted policy

The arrows show the greedy action selected by the best-performing Q-table. **S** is the start, **C** marks the cliff, and **G** is the goal.

```
In [5]: best_episodes = int(summary_df.loc[summary_df["Test success rate (%)"].idxmax(), "T
policy = np.argmax(results[best_episodes]["q_table"], axis=1).reshape(4, 12)
symbols = {0: "↑", 1: "→", 2: "↓", 3: "←"}
cliff = np.array([list(row) for row in ["FFFFFFFFFFFF", "FFFFFFFFFFFF", "FFFFFFFFFFFF

fig, axes = plt.subplots(1, 2, figsize=(10, 4))
env = gym.make("CliffWalking-v1", render_mode="rgb_array")
env.reset(seed=SEED)
axes[0].imshow(env.render())
axes[0].set_title("CliffWalking environment")
axes[0].axis("off")
env.close()

axes[1].imshow(np.zeros((4, 12)), cmap="Blues", vmin=0, vmax=1)
for row in range(4):
    for col in range(12):
        label = cliff[row, col] if cliff[row, col] in "SCG" else symbols[policy[row, col]]
        axes[1].text(col, row, label, ha="center", va="center", fontsize=20)
axes[1].set(xticks=range(12), yticks=range(4), xticklabels=[], yticklabels=[], title="Greedy policy (500 episodes)")
axes[1].grid(color="white", linewidth=2)
plt.tight_layout()
plt.show()
print("Policy action numbers (0=up, 1=right, 2=down, 3=left):")
print(policy)
```



Policy action numbers (0=up, 1=right, 2=down, 3=left):

```
[[1 1 1 1 1 1 1 1 1 1 1 2]
 [1 1 1 1 1 1 1 1 1 1 1 2]
 [1 1 1 1 1 1 1 1 1 1 1 2]
 [0 0 0 0 0 0 0 0 0 0 0 0]]
```

```
In [6]: # Minimal self-check for the tabular experiment.
assert set(summary_df["Training episodes"]) == set(EPISEODE_CONFIGS)
assert all(len(result["rewards"]) == episodes for episodes, result in results.items)
assert summary_df["Test success rate (%)"].between(0, 100).all()
print("Lab checks passed.")
```

Lab checks passed.